

Everyone knows **what** the news said. Nobody profiles **how** the sources behave.


Some questions you might want answered about who gives you the news:

- ▶ How mainstream is any given news source?
- ▶ Which sources break real stories first — unique, and ahead of the pack?
- ▶ Which just echo consensus after it's already been published elsewhere?

Narrative Nexus is an instrument that answers them.

For anyone whose money or decisions ride on news — hedge funds, PR firms, risk analysts.

37 outlets, scored by behavior — not by opinion.



Every dot is a news outlet, placed by two measured behaviors: how often it **breaks claims early** (x) vs. how often those claims **survive into consensus** (y).

Four archetypes emerge:

Early Breaker

Unmatched Breaker

Late but Reliable

Consensus Echo


No composite score. No editorial judgment.

"Narrative Nexus tracks consensus reality, not truth."

Four AI agents. All inference on AMD Instinct via Fireworks AI.

- ▶ **1 Intake & Clustering** — embeds articles, groups them into story clusters
- ▶ **2 Forensic Extraction** — strips spin, extracts atomic factual claims (LLM)
- ▶ **3 Consensus Alignment** — pure math: where do ≥ 2 independent sources converge?
- ▶ **4 Silent Auditor** — re-reads old articles, flags unannounced edits

All AI stages are configured to run on Fireworks AI, serving inference on AMD Instinct MI300X/MI250X accelerators. The shipped database was built end-to-end through Fireworks during hackathon week using AMD hackathon credits. Provider-agnostic by design: every stage's provider and model is independently configurable —



One pipeline, run end-to-end on two real stories.

Article scraping was performed July 3–5, 2026; coverage windows below reflect when the news itself was published.

US-Iran War: March Escalation & April Ceasefire

48-day arc

Mar 9, 2026 – Apr 26, 2026

6 articles 19 claims 3 sources 1 absorbed

What consensus formed on

On Tuesday, the U.S. and Israel launched airstrikes against Iran.

Reported by

reuters.com, theguardian.com, apnews.com

Time to consensus

Median: 117.8 days

From first report to consensus-absorbed

[View Cluster Report →](#)

[🕒 View Timeline](#)

Venezuela Emergency and Rescue Response

5-day surge

Jun 24, 2026 – Jun 29, 2026

61 articles 138 claims 20 sources 3 absorbed

What consensus formed on

On Thursday, June 25, 2026, Interim President Delcy Rodríguez reported on state-run TV that at least 164 people were killed in the earthquakes.

The earthquakes caused damage in Caracas and other areas.

Rescue operations are underway in Venezuela.

Reported by

apnews.com, bbc.com, npr.org, theguardian.com, foxnews.com, nytimes.com, www3.nhk.or.jp, france24.com, theintercept.com, zerohedge.com...

Silent edits detected

10 edits detected

Time to consensus

Median: 10.9 days

From first report to consensus-absorbed

[View Cluster Report →](#)

[🕒 View Timeline](#)

US-Iran War — a 48-day arc: one outlet's scoop ripening into cross-source consensus, animated claim by claim on the Timeline.

Venezuela Emergency — a 5-day surge: 20 outlets, 138 claims, 10 silent edits detected. Corroborated facts vs. single-outlet outliers — Maduro's arrest reported by one source, India's field hospital only by The Hindu, China's aid only by Global Times.

Consensus requires ≥ 2 independent sources. **Single-source claims can never self-validate.**

Collecting news is itself the first finding.

Here is what we found just trying to get the news online at all:

- ▶ **Half the news is gated.** Of 358 articles collected, only 182 yielded full body text — the rest paywalled or truncated RSS. Three major outlets started at 0% extraction.
- ▶ **A handful of wires anchor everything.** Consensus reality is effectively defined by ~5 wire services and public broadcasters; most outlets echo downstream.
- ▶ **Most stories have one voice.** 9 of 17 story clusters were covered by a single source. Corroboration is rarer than it looks.
- ▶ **Global coverage doesn't come free.** First panel draft: 90% Global North. We empirically tested 45 candidate feeds to reach 37 sources, 7 regions, 6 continents.

Our first engine reported 2,625 consensus claims. 96.8% were an artifact.

A claim from a single source, in a cluster with no other sources, computed 1-of-1 = 100% agreement → "consensus." **The machine was agreeing with itself.**

The fix: consensus now requires **≥2 independent consensus-pool sources** reporting the same claim, matched semantically across outlets (cosine ≥ 0.85 — A/B tested; 0.80 introduced factually wrong merges).

Every threshold in the pipeline is empirically calibrated — clustering eps swept against hand-labeled story groups, cluster-size guards against cross-story blobs, per-vertical consensus thresholds.

2,625



10

Smaller number. **Real number.**

Built for scrutiny.

- ▶ **No composite score, ever** — six independent behavioral dimensions; the user decides what matters.
- ▶ **"Consensus reality, not truth"** — the system never says a source was right or wrong. That constraint is what makes it defensible and non-partisan.
- ▶ **Nothing mocked** — the shipped database is one verified pipeline run: 358 articles, 378 claims, 13,653 daily reputation snapshots. Every number on every page traces to a query.
- ▶ **Live by design** — the same pipeline can collect, cluster, extract, and score continuously; press Start in Settings. Each clone is its own instance.

THE INSTRUMENT

The product isn't the run. It's the instrument.

What you've seen is one verified pass over a real news corpus — 358 articles, 37 sources, 6 continents, processed by the real pipeline. Nothing mocked. **This run was short by design** — enough to prove the pipeline works end to end.

The real product is the instrument itself: as far as we've found, **nobody is measuring how news outlets behave over time the way this does**. Run it continuously and that behavioral record only gets richer.

An instrument for telling the clones from the outliers.