

NARRATIVE NEXUS

We all know news sources carry biases — left, right, or somewhere harder to place. And we all have our favorite outlets. But how much do you really know about the ongoing tendencies of any one source?

Some questions you might want answered about who gives you the news:

- How mainstream is any given news source?
- Which sources break real stories first — unique, and ahead of the pack?
- Which just echo consensus after it's already been published elsewhere?

Narrative Nexus is an instrument built to answer them. It doesn't score bias — it scores behavior, measured over time. Of particular interest to anyone whose money rides on the news: hedge funds, PR firms, risk analysts.

Constraint: Narrative Nexus tracks consensus reality, not truth.

THE INSTRUMENT

37 outlets, scored by behavior — not by opinion.


Every dot is a news outlet, placed by two measured behaviors:

- How often it breaks claims early (x-axis)
- How often those claims survive into consensus (y-axis)

Four archetypes emerge from those two axes — **Early Breaker**, **Unmatched Breaker**, **Late but Reliable**, **Consensus Echo** — read straight off the quadrants.

There is no composite score, and no editorial judgment.

"Narrative Nexus tracks consensus reality, not truth."



Sources reputation map

THE MACHINE

Four agents – two AI powered, two deterministic. All AI inference on **AMD** Instinct via Fireworks AI.


All AI stages are configured to run on Fireworks AI, serving inference on **AMD** Instinct MI300X/MI250X accelerators. The shipped database was built end-to-end through Fireworks during hackathon week using **AMD** hackathon credits.

Provider-agnostic by design: every stage's provider and model is independently configurable.

Gemma integration: **Gemma** 4 E4B (on-demand via Fireworks) is a selectable provider, verified running our claim-extraction prompt across the full 61-article Venezuela cluster — 268 structured claims extracted. Evidence: docs/evidence/gemma/ in the repo.

The Pipeline Stages:

- Intake & Clustering — Embeds articles, groups them into story clusters.
- Forensic Extraction — Strips spin, extracts atomic factual claims via LLMs.
- Consensus Alignment — Pure mathematics: where do ≥ 2 independent sources converge?
- Silent Auditor — Re-reads old articles, flags unannounced post-publication edits (deterministic text diff).




Pipeline flow

THE PROOF

One pipeline, run end-to-end on two real stories.

- US-Iran War (a 48-day arc): one outlet's scoop ripening into cross-source consensus, animated claim by claim on the Timeline.



US-Iran timeline

THE PROOF

- Venezuela Emergency (a 5-day surge): 20 outlets, 138 claims, 5 silent edits detected — corroborated facts vs. single-outlet outliers (Maduro's arrest by one source; India's field hospital only in The Hindu; China's aid only in Global Times).

Consensus requires ≥ 2 independent sources. Single-source claims can never self-validate.



Venezuela cluster report

WHAT YOU DON'T SEE

Collecting news is itself the first finding. Here is what we found just trying to get the news online at all:

- Half the news is gated. Of 358 articles collected, only 182 yielded full body text — the rest were paywalled or truncated RSS. Three major outlets started at 0% extraction.
- A handful of wires anchor everything. Consensus reality is effectively defined by ~5 wire services and public broadcasters; most outlets echo downstream.
- Most stories have one voice. 9 of 17 story clusters were covered by a single panel source — partly because regional and contrarian outlets cover stories nobody else touches. Corroboration is rarer than it looks.
- Global coverage doesn't come free. First panel draft was 90% Global North. We empirically tested 45 candidate feeds to reach 37 sources across 7 regions and 6 continents.

The implication: across this panel, most events are carried by a single outlet — **independent corroboration is the exception, not the rule.**

MORE FINDINGS...

Our first engine reported 2,625 consensus claims. 96.8% were an artifact.

A claim from a single source, in a cluster with no other sources, computed 1-of-1 = 100% agreement → "consensus." The machine was agreeing with itself.

The Fix: Consensus now requires ≥ 2 independent consensus-pool sources reporting the same claim, matched semantically across outlets using cosine similarity (≥ 0.85 — A/B tested; 0.80 introduced factually wrong merges).

Every threshold in the pipeline is empirically calibrated:

- Clustering eps swept against hand-labeled story groups.
- Cluster-size guards applied against cross-story blobs.
- Custom per-vertical consensus thresholds enforced.

2,625 claims (artifact) → 10 claims (real, filtered number)

DISCIPLINE

Built for scrutiny.

- No composite score, ever: Six independent behavioral dimensions are tracked; the user decides what matters.
- "Consensus reality, not truth": the system never rules a claim right or wrong — an uncorroborated claim is unabsorbed, not false. A deliberate design choice: it's what makes the instrument defensible and non-partisan.
- Nothing mocked: The shipped database is one verified pipeline run containing 358 articles, 378 claims, and 13,653 daily reputation snapshots. Every number on every page traces directly to a database query.

What we deliberately can't claim (yet): we set out to spot the "honest, brave" source — the outlet that breaks a true claim no one else will touch. But an instrument that measures consensus, by construction, can't tell a lone source that was right from one that was wrong: both stay unabsorbed until corroboration arrives. That's not a gap to patch — it's the same fact as the moat. Only a longitudinal record, run long enough for time to settle it, could ever answer that.


THE INSTRUMENT

The product isn't the run. It's the instrument.

What you've seen is one verified pass over a real news corpus — 358 articles, 37 sources, 6 continents, processed by the real pipeline. Nothing mocked. This run was short by design — enough to prove the pipeline works end-to-end.

The real product is the instrument itself: as far as we've found, nobody is measuring how news outlets behave over time the way this does. Run it continuously and that behavioral record only gets richer. The moat is time.

An instrument for telling the clones from the outliers.



Stories page